

Lab 3: Using Classification to Build a Marketing Target Model

Umais Siddiqui

2026-02-14

Table of contents

1	Problem 1: What Can Machine Learning Do?	2
1.1	(a) How ML Can Help the Marketing Team	2
1.2	(b) The Kaggle Notebook Flaw	2
2	Problem 2: EDA and Baseline Model	3
2.1	(a) EDA and Temporal Patterns	3
2.1.1	Dataset Overview	3
2.1.2	Missing Values	3
2.1.3	Temporal Success Rate Analysis	4
2.1.4	Success Rate by Year	5
2.1.5	Key Predictor Relationships	5
2.1.6	Contact Method and Population Composition Over Time	5
2.1.7	Summary of Key Findings from EDA	6
2.2	(b) Baseline Model and Data Drift	6
2.2.1	Train-Test Split	6
2.2.2	Define Epochs Based on Macroeconomic Conditions	7
2.2.3	Why Temporal Variation Causes Metric Mismatch	7
2.2.4	Baseline Logistic Regression Model	7
2.2.5	ROC Curves by Epoch	8
2.2.6	Baseline Model Performance by Epoch	8
2.3	(c) Lift Curves and Hit Rates	9
2.3.1	Lift and Hit Rate Summary	9
3	Problem 3: Model Interpretation	10
3.1	(a) Odds Ratios	10
3.1.1	Top 10 Positive Predictors	10
3.1.2	Top 10 Negative Predictors	10
3.2	(b) Marginal Effects	12
4	Problem 4: Model Comparison and Refinement	13
4.1	(a) Refining Logistic Regression	13
4.1.1	Feature Engineering	13

4.1.2	Refined Model Training	13
4.1.3	Lift Curves: Baseline vs Refined	13
4.1.4	Refined Model Performance	14
4.2	(b) Naive Bayes Comparison	14
4.2.1	ROC Curves: Logistic Regression vs Naive Bayes	14
4.2.2	Model Comparison Results	15
5	Extra Credit: How Has the Marketing Team Already Been Optimizing?	16
5.0.1	Success Rate by Previous Contacts	16
5.0.2	Mean Duration of Successful Calls by Year	16
5.0.3	Job Distribution by Year	16
5.1	Evidence of Marketing Team Optimization	17
6	Conclusion	18

1 Problem 1: What Can Machine Learning Do?

1.1 (a) How ML Can Help the Marketing Team

A classification model trained on historical campaign data can help the marketing team in several concrete ways. First, and most directly, it can score and rank prospective clients by their predicted probability of subscribing to a term deposit, allowing the team to prioritize high-probability leads and avoid wasting calls on unlikely converters. This targeted approach increases the hit rate (fraction of successful calls) and reduces the cost per acquisition. If the model can identify the top 10% of clients who are, say, 3× more likely to subscribe, the team could achieve the same number of conversions with far fewer calls — or achieve many more conversions with the same budget.

Second, the model can provide interpretable insights into which client characteristics and campaign strategies drive success. For example, if the model reveals that cellular contact is far more effective than telephone, or that clients who previously subscribed are the easiest to re-convert, the team can redesign their outreach strategy accordingly. The model can also help identify the best timing for campaigns by revealing how macroeconomic conditions (like the Euribor interest rate) affect client receptivity — the team might accelerate campaigns when conditions are favorable and scale back when the environment is hostile to term deposit sales.

1.2 (b) The Kaggle Notebook Flaw

The most upvoted Kaggle notebook includes `duration` (the length of the phone call in seconds) as a predictor variable. This is a fatal flaw because duration is only known **after** the call has been made — it cannot be used to decide who to call. As the UCI documentation explicitly warns, `duration` should be discarded for any realistic predictive model. Including it inflates model performance dramatically (since longer calls strongly predict success), but the model is completely useless in practice because you would need to already have called the client to know how long the call lasted.

Table 1: Dataset Overview

	Metric	Value
Total	Observations	45,211
	Features	26
	Success Rate	11.70%
	Successes	5,289
	Failures	39,922

2 Problem 2: EDA and Baseline Model

2.1 (a) EDA and Temporal Patterns

2.1.1 Dataset Overview

The dataset contains 45,211 observations with 24 features. The target variable y indicates whether a client subscribed to a term deposit, with a baseline success rate of approximately 11.7%.

2.1.2 Missing Values

No missing values in any column.

Table 2: Success Rate by Year with Macroeconomic Indicators

	n	success_rate	mean_euribor	mean_emp_var_rate
year				
2008	27729	0.0505	4.8861	1.1142
2009	14862	0.1706	1.6231	-2.0457
2010	2620	0.5160	0.7657	-1.6120

2.1.3 Temporal Success Rate Analysis




Figure 1: Temporal patterns in campaign success rate and volume

2.1.4 Success Rate by Year

2.1.5 Key Predictor Relationships




Figure 2: Key predictor relationships with the target variable

2.1.6 Contact Method and Population Composition Over Time




Figure 3: Contact method and population composition over time

2.1.7 Summary of Key Findings from EDA

The most striking pattern in this dataset is the **dramatic temporal variation** in campaign success. The success rate rises from roughly 5% in 2008 to 17% in 2009 to over 50% in 2010. This variation is primarily driven by macroeconomic conditions, specifically the Euribor 3-month interest rate and the employment variation rate.

The **Euribor 3-month rate** — the interbank lending rate in the Eurozone — declined sharply from ~5% in mid-2008 to below 1% by late 2009, driven by the ECB’s aggressive rate cuts during the 2008–2009 financial crisis. When market interest rates are low, bank term deposits become relatively more attractive (they offer fixed returns in a low-rate environment) and the opportunity cost of locking up money is lower. The correlation between the `euribor3m` rate and campaign success is strongly negative: as the rate fell, success rates soared.

The **employment variation rate** (`emp.var.rate`) captures labor market conditions. It turned sharply negative during the crisis (from +1.4 to -3.4), reflecting rising unemployment. Counterintuitively, worse employment conditions are associated with higher term deposit uptake — likely because economic uncertainty makes safe, fixed-return savings products more appealing to risk-averse consumers.

Beyond macroeconomics, the composition of the target population changed over time. In 2010, the share of retired people and students called was notably higher — these are the highest-success-rate job categories. The contact method also shifted: the first two months (May–June 2008) used only “unknown” contact, while later periods predominantly used cellular.

Other important predictors include: `poutcome` (clients who previously succeeded are ~65% likely to subscribe again), `housing` (those without housing loans subscribe at higher rates), `contact` (cellular outperforms telephone), and `education` (tertiary-educated clients subscribe more).

2.2 (b) Baseline Model and Data Drift

2.2.1 Train-Test Split

Training set: 36,168 observations

Testing set: 9,043 observations

Train success rate: 0.1170

Test success rate: 0.1170

Stratification choice: I stratify by `y` (the target variable) rather than by time. The target is imbalanced (~11.7% positive), so stratification ensures both train and test sets preserve this ratio. I do not stratify by time because the assignment specifically asks us to examine how model performance varies across temporal epochs — a random (non-temporal) split serves this purpose, as each epoch will be represented in both train and test sets.

2.2.2 Define Epochs Based on Macroeconomic Conditions

Test set epoch distribution:

euribor3m

Low (<1.5%) 2679

Medium (1.5-3.5%) 777

High (>3.5%) 5587

Test set success rate by epoch:

Low (<1.5%): 0.2628 (n=2679)

Medium (1.5-3.5%): 0.1081 (n=777)

High (>3.5%): 0.0483 (n=5587)

2.2.3 Why Temporal Variation Causes Metric Mismatch

If we train and evaluate a model on a random split of the entire dataset, the overall performance metrics blend together periods with very different success rates and different predictor distributions. In 2008, only 5% of calls succeeded (high euribor, pre-crisis); in 2010, over 50% succeeded (low euribor, post-crisis). A model evaluated on the pooled test set might show decent accuracy — but this average hides the fact that the model may perform brilliantly in one epoch and terribly in another.

In practice, the marketing team uses the model in real time, where they face one macroeconomic regime at a time. If the current conditions resemble 2008, the pooled metrics will be misleadingly optimistic. Conversely, if conditions resemble 2010, the pooled metrics may understate performance. Evaluating by epoch reveals these regime-specific performance characteristics, which is exactly what the marketing team needs to set realistic expectations.

2.2.4 Baseline Logistic Regression Model

Baseline logistic regression fitted.

Table 3: Baseline Model Performance by Epoch

	Epoch	N	Success Rate	Accuracy	Precision	Recall	F1	AUC
	Low (<1.5%)	2679	0.2628	0.7850	0.6649	0.3665	0.4725	0.7859
	Medium (1.5–3.5%)	777	0.1081	0.8919	0.0000	0.0000	0.0000	0.6777
	High (>3.5%)	5587	0.0483	0.9517	0.0000	0.0000	0.0000	0.5967
	Overall	9043	0.1170	0.8972	0.6649	0.2439	0.3568	0.7956

2.2.5 ROC Curves by Epoch




Figure 4: ROC curves by epoch for the baseline logistic regression model

2.2.6 Baseline Model Performance by Epoch

Observations on epoch-specific performance:

The model performance varies dramatically across epochs. In the **Low euribor epoch** (post-crisis, high success rate $\sim 26\%$), the model achieves a reasonable AUC of ~ 0.79 and meaningful precision/recall tradeoffs. In the **Medium epoch** (transition, $\sim 11\%$ success), the AUC drops to ~ 0.68 and the model struggles — precision and recall at the default threshold collapse because the model was trained on a mix of all epochs and the signal is weaker here. In the **High euribor epoch** (pre-crisis, only $\sim 5\%$ success), the AUC is lowest (~ 0.60) and the model essentially predicts “no” for everyone, achieving high accuracy (95%) purely from the class imbalance but with zero recall.

The overall test metrics are misleading: the AUC of ~ 0.80 looks respectable, but this is driven almost entirely by the model’s ability to separate the easy Low-epoch observations from the rest. A marketer working in 2008-like conditions would find this model far less useful than the overall metrics suggest. This confirms the need for epoch-stratified evaluation.

Table 4: Baseline Model: Lift and Hit Rate Summary

	Epoch Base Rate	Lift @ 10%	Hit Rate @ 10%	Lift @ 50%	Hit Rate @ 50%
Low (<1.5%)	0.2628	2.78	0.7303	1.65	0.4324
Medium (1.5–3.5%)	0.1081	2.64	0.2857	1.48	0.1598
High (>3.5%)	0.0483	1.56	0.0753	1.26	0.0609

2.3 (c) Lift Curves and Hit Rates




Figure 5: Lift curves by epoch for the baseline model

2.3.1 Lift and Hit Rate Summary

Interpretation of lift curves:

The lift curves reveal the model’s practical value for campaign targeting. In the **Low euribor epoch** (favorable conditions), the model achieves a lift of ~2.8 at 10% targeting — meaning if the team calls only the top-ranked 10% of clients, they will find subscribers at $2.8\times$ the base rate, achieving a hit rate of ~73%. Even at 50% targeting, the lift is ~1.65 with a hit rate of ~43%. This is genuinely useful: the team can cut their call volume in half while capturing most of the conversions.

In the **Medium epoch**, the model still provides value with a lift of ~2.6 at 10%, but the hit rate is only ~29% because the underlying base rate is much lower (~11%). In the **High euribor epoch** (hostile conditions, 5% base rate), the lift at 10% is only ~1.6 with a hit rate of ~8%. The model provides some targeting value, but the absolute hit rates are low because there are simply fewer willing clients in the population during this period.

This analysis demonstrates that **lift is more stable across epochs than raw accuracy metrics**, and it directly translates to ROI: a lift of $2\times$ means the team gets twice the return per call compared to random dialing.

3 Problem 3: Model Interpretation

3.1 (a) Odds Ratios




Figure 6: Odds ratios for the baseline logistic regression model (top 10 positive and top 10 negative)

Note: Numeric features are standardized, so odds ratios reflect the effect of a 1-SD increase in each numeric predictor.

3.1.1 Top 10 Positive Predictors

3.1.2 Top 10 Negative Predictors

Interpretation of odds ratios:

Because the numeric features are standardized, their odds ratios represent the multiplicative change in odds per one standard deviation increase. Key findings:

Table 5: Top 10 Positive Predictors (Highest Odds Ratios)

Feature	Coefficient	Odds Ratio
poutcome_success	1.791856	6.000579
cons.price.idx	0.743086	2.102414
poutcome_unknown	0.606164	1.833384
cons.conf.idx	0.297463	1.346439
education_tertiary	0.255110	1.290603
poutcome_other	0.246305	1.279289
job_retired	0.212605	1.236896
nr.employed	0.152375	1.164597
education_secondary	0.151417	1.163482
job_unemployed	0.092521	1.096936

Table 6: Top 10 Negative Predictors (Lowest Odds Ratios)

Feature	Coefficient	Odds Ratio
job_unknown	-0.133637	0.874908
campaign	-0.137486	0.871546
contact_telephone	-0.209296	0.811155
loan_yes	-0.225847	0.797840
marital_married	-0.274965	0.759599
job_housemaid	-0.305227	0.736956
housing_yes	-0.404195	0.667514
euribor3m	-0.428069	0.651766
emp.var.rate	-0.868164	0.419721
contact_unknown	-1.187916	0.304856

Table 7: Average Marginal Effect of Contact Method (Baseline: telephone)

Average Marginal Effect of Contact Method
(Baseline: telephone)

	Epoch	AME: Cellular vs Telephone	AME: Unknown vs Telephone	Mean P(telephone)	Mean I
Low (<1.5%)		+0.0312		-	
0.1149	0.2407	0.2719	0.1258		
Medium (1.5–3.5%)		+0.0163		-	
0.0475	0.0791	0.0954	0.0316		
High (>3.5%)		+0.0149		-	
0.0430	0.0713	0.0862	0.0283		

- **Largest positive effects:** The strongest predictor of success is `poutcome_success` (OR = 6.0) — clients who subscribed in a previous campaign are about 6× more likely to subscribe again. This is the most actionable finding: re-targeting previous subscribers is extremely effective. The `cons.price.idx` (OR = 2.1) reflects macroeconomic favorability, and `poutcome_unknown` (OR = 1.8) captures that first-time contacts have a higher baseline than clients who previously refused.
- **Largest negative effects:** The most harmful predictor is `contact_unknown` (OR = 0.30) — not knowing the contact method is associated with 70% lower odds, likely because “unknown” contact dominated the earliest, least effective months. The macroeconomic indicators `emp.var.rate` (OR = 0.42) and `euribor3m` (OR = 0.65) confirm that worsening economic conditions (higher employment variation rate, higher interbank rates) strongly reduce subscription likelihood. Having a housing loan (OR = 0.67) reduces odds by about a third, likely because clients with existing loan obligations are less willing to lock up additional capital.

3.2 (b) Marginal Effects

Interpretation of marginal effects:

The average marginal effects reveal how switching contact method changes the predicted probability of success, holding all other client characteristics constant. Across all epochs, **cellular contact is associated with a modest improvement over telephone**: the AME ranges from about +1.5 percentage points in the High epoch to +3.1 percentage points in the Low epoch. This suggests that reaching clients on their cell phone is slightly more effective, possibly because it catches them at more convenient moments or signals a more personal touch.

The **unknown contact method is dramatically worse than telephone**, with AMEs ranging from -4.3 to -11.5 percentage points. This large negative effect likely reflects the fact that “unknown” contact was used almost exclusively in May–June 2008, the earliest and least successful period

of the campaign. The “unknown” contact category may thus be partially confounded with other early-campaign characteristics rather than representing a pure contact method effect.

The marginal effects are larger in magnitude in the **Low euribor epoch** because the baseline success probability is higher there — the logistic function is steeper near $P=0.3$ than near $P=0.05$, so the same coefficient translates to a larger probability shift.

4 Problem 4: Model Comparison and Refinement

4.1 (a) Refining Logistic Regression

4.1.1 Feature Engineering

Engineered features:

- age_sq: Captures non-linear age effect (young and old subscribe more)
- log_campaign: Diminishing returns of repeated contact attempts
- log_balance: Log-transform of balance (right-skewed)
- contacted_before: Binary flag for whether client was previously contacted
- prev_success: Binary flag for previous campaign success
- is_retired_or_student: Flags the two highest-success job categories

4.1.2 Refined Model Training

Refined logistic regression fitted.

Regularization: $C=0.5$ (slightly more regularization than default)

Features: 14 numeric + 8 categorical

4.1.3 Lift Curves: Baseline vs Refined



Figure 7: Lift curves: Baseline vs Refined logistic regression by epoch

Table 8: Refined Model: Lift and Hit Rate Summary

	Epoch	Base Rate	Lift @ 10%	Hit Rate @ 10%	Lift @ 50%	Hit Rate @ 50%
Low	(<1.5%)	0.2628	2.78	0.7303	1.64	0.4317
Medium	(1.5–3.5%)	0.1081	2.28	0.2468	1.45	0.1572
High	(>3.5%)	0.0483	1.78	0.0860	1.29	0.0623

4.1.4 Refined Model Performance

Comparison to baseline:

The refined model incorporates several improvements: (1) `age_sq` captures the non-linear age effect (young students and older retirees subscribe at higher rates than middle-aged workers), (2) `log_campaign` accounts for diminishing returns of repeated contacts, (3) `log_balance` better handles the skewed balance distribution, (4) explicit binary features for `contacted_before` and `prev_success` provide cleaner signal than the raw `pdays/previous` variables, and (5) `is_retired_or_student` directly encodes the highest-success job categories. The regularization parameter was set to $C=0.5$ to reduce overfitting given the increased feature count.

The refined model shows modest improvements in the **High euribor epoch**, where the additional features help discriminate in a low-signal environment. In the **Low epoch**, performance is similar — there is already strong signal from the macroeconomic variables. The improvements are incremental rather than dramatic, which is expected: the biggest gains come from including the macroeconomic indicators (already in the baseline), and the individual-level features add smaller marginal value.

4.2 (b) Naive Bayes Comparison

Gaussian Naive Bayes fitted.

4.2.1 ROC Curves: Logistic Regression vs Naive Bayes




Figure 8: ROC curves: Logistic Regression vs Naive Bayes by epoch

Table 9: Model Comparison: Logistic Regression vs Naive Bayes

	Epoch	LR AUC	NB AUC	LR Lift@10%	NB Lift@10%	LR Hit@10%	NB Hit@10%	LR Lift@50%	NB Lift@50%
Low (<1.5%)		0.7859	0.7014	2.78	2.62	0.7303	0.6891	1.65	1.65
Medium (1.5–3.5%)		0.6777	0.5676	2.64	1.56	0.2857	0.1688	1.48	1.48
High (>3.5%)		0.5967	0.5765	1.56	1.63	0.0753	0.0789	1.26	1.26

4.2.2 Model Comparison Results

Naive Bayes vs Logistic Regression:

Gaussian Naive Bayes consistently underperforms logistic regression across all epochs, with lower AUC, lower lift, and lower hit rates.

The fundamental issue is that Naive Bayes assumes all features are **conditionally independent** given the class label — an assumption that is clearly violated here. The macroeconomic indicators (`euribor3m`, `emp.var.rate`, `nr.employed`, `cons.price.idx`, `cons.conf.idx`) are all highly correlated with each other (they all track the business cycle), and treating them as independent leads to overconfident, poorly calibrated probability estimates. Additionally, applying Gaussian assumptions to the one-hot-encoded categorical features is a rough approximation at best — these binary features are not normally distributed.

Logistic regression handles correlated features naturally through its linear combination, and makes no distributional assumptions about the inputs. For this dataset with strong feature correlations and a mix of continuous and categorical predictors, **logistic regression is clearly the better choice**. I was not able to outperform logistic regression with Naive Bayes in any epoch.

5 Extra Credit: How Has the Marketing Team Already Been Optimizing?




Figure 9: Evidence of marketing team optimization over time

5.0.1 Success Rate by Previous Contacts

5.0.2 Mean Duration of Successful Calls by Year

Mean duration of SUCCESSFUL calls by year:

2008: 838.4 seconds (n=1401)

2009: 459.7 seconds (n=2536)

2010: 370.9 seconds (n=1352)

5.0.3 Job Distribution by Year

Table 10: Success rate by number of previous contacts

	n	success_rate
previous		
0	36954	0.0916
1	2772	0.2103
2	2106	0.2165
3	1142	0.2574
4	714	0.2395
5	459	0.2636
6	277	0.2996
7	205	0.2634

Job distribution by year:

2008 (n=27729):

blue-collar: 22.9%
management: 21.1%
technician: 18.1%
admin.: 10.4%

2009 (n=14862):

blue-collar: 21.4%
management: 20.0%
technician: 14.7%
admin.: 13.0%

2010 (n=2620):

management: 23.5%
retired: 15.0%
technician: 14.5%
admin.: 14.0%

5.1 Evidence of Marketing Team Optimization

Several patterns in the data suggest that the marketing team was actively learning and adapting their strategy over time, beyond what macroeconomic conditions alone would explain:

1. **Shifting to higher-success demographic segments.** The share of retired and student clients in the target list increased dramatically from ~5% in 2008 to ~24% in 2010. These are the two job categories with the highest subscription rates. This suggests the team identified

these segments as high-value and shifted their targeting accordingly — a clear exploitation strategy.

2. **Reducing campaign intensity (fewer calls per person).** In mid-2008, the team was making an average of 3–4 contacts per client (with some receiving 50+ calls). By 2009–2010, this dropped to ~1.7–2 contacts. The team appears to have learned that repeatedly calling reluctant clients has diminishing returns and generates ill will, and shifted toward making fewer, more targeted calls.
3. **Adopting cellular contact.** The campaign shifted from 100% “unknown” contact in May–June 2008 to predominantly cellular contact from July 2008 onward. Cellular contact has a higher success rate (as we saw in the marginal effects analysis). This shift could reflect learning about the superior effectiveness of reaching clients on their mobile phones.
4. **Shrinking volume with increasing quality.** The call volume declined dramatically from ~28,000 calls in 2008 to ~2,600 in 2010, while the success rate increased from 5% to 52%. While much of the success rate increase is macroeconomic, the massive volume reduction suggests the team was becoming more selective — the classic exploitation side of the explore-exploit tradeoff. In 2008, they cast a wide net (exploration), gathering data on what works. By 2010, they were making fewer, better-targeted calls (exploitation).
5. **Selection bias in the sample.** Because the team was increasingly targeting high-success segments, the dataset is not a random sample of the bank’s customers. Segments that the team learned were unlikely to subscribe (e.g., blue-collar workers with housing loans during high-rate periods) are underrepresented in later time periods. This means any model trained on this data inherits the team’s implicit targeting model — it cannot discover opportunities in populations that the team already decided not to call. This is the fundamental explore-exploit tradeoff: the team’s exploitation strategy improves short-term results but limits the data available for model improvement.

6 Conclusion

This analysis demonstrates several key insights for building and deploying classification models in marketing contexts:

1. **Temporal variation matters.** Macroeconomic conditions (particularly the Euribor 3-month rate) are the strongest predictors of campaign success, creating distinct “epochs” with dramatically different baseline success rates. Models must be evaluated within these epochs to provide actionable guidance.

2. **Lift curves are more informative than accuracy.** In imbalanced classification problems, lift curves directly translate to business value — a $2\times$ lift means twice the ROI per call compared to random targeting.
3. **Feature engineering provides incremental gains.** While macroeconomic indicators dominate, engineered features like `age_sq`, `log_campaign`, and `is_retired_or_student` provide modest improvements, especially in low-signal epochs.
4. **Model selection depends on data characteristics.** Naive Bayes fails here due to violated independence assumptions among correlated macroeconomic indicators. Logistic regression's flexibility with correlated features makes it the better choice.
5. **Real-world optimization affects model training data.** The marketing team's adaptive targeting strategy creates selection bias in the historical data, limiting the model's ability to discover opportunities in under-sampled segments.